

Phase 1: Brainstorming & Ideation

AI Comic Story Creator — SmartBridge Project Submission

1. Problem Statement

Creating a comic book traditionally requires a writer, an artist, and a letterer working in sequence. This project explores whether a multi-agent AI pipeline (Google Gemini for writing/direction + Imagen 3 for illustration) can collapse this into a single automated system that turns a short text prompt into a finished, multi-panel comic.

2. Target Users

- Hobbyist storytellers with no drawing skill who want to visualize their stories.
- Indie creators prototyping issue concepts quickly before committing to full production.
- Educators building comic-strip explainers for classroom material.
- Marketing teams producing short branded mini-comics.

3. Core Idea

A decoupled, agentic pipeline where Gemini acts as the **Director & Writer** and Imagen acts as the **Illustrator**: Story Architect → Character Consistency Manager → Panel Layout Director → Imagen Rendering → Typesetting/Compositing → exported PDF/CBZ.

4. Possible Differentiators

- Character consistency across panels via a reusable 'Visual Anchor String' injected into every image prompt.
- Style-switching: the same script can be rendered in multiple art styles (Manga, Noir, Silver Age, Cyberpunk).
- Automatic dialogue-bubble placement instead of manual lettering.
- Flexible export: print-ready 300 DPI CMYK PDF, digital CBZ/CBR, or Webtoon-style vertical scroll.

5. Stretch Feature Ideas

- User-editable script step before rendering, so users can tweak beats before spending image-generation credits.
- 'Regenerate this panel only' instead of re-rendering a full page.
- Genre presets (sci-fi, romance, horror) that pre-load tone and pacing templates for the Story Architect agent.
- Optional audio narration layer for accessibility or animated-comic export.

6. Risks Identified Early

- API cost: multiple Imagen calls per page across a 7-8 page issue can add up quickly.
- Character drift even with visual anchors, since this is prompt-based consistency, not model fine-tuning.

- Generation latency requiring an async job/queue design and a 'still generating' UX state.
- Content moderation for generated art and dialogue text.